



Kubernetes The CrashLoopBackOff Periodic Table




A Definitive Guide to Decoding &
Resolving Kubernetes Container
Failures



What is CrashLoopBackOff?

It is not an error itself, but a state. It indicates that a container is crashing repeatedly, and Kubernetes is implementing an exponential backoff delay (10s, 20s, 40s...) before trying again to protect the node's resources.

CrashLoopBackOff Error in Kubernetes





1.Exit Code 1 – Application Error

The "Why":

The most common "catch-all" error. The application started but encountered a fatal runtime exception (e.g., a missing semicolon in a config file, a null pointer, or a failed dependency).

The Deep Dive:

Check the standard output. Usually, the application framework (Node, Python, Java) will dump a stack trace here.

The Fix:

- Run `kubectl logs <pod-name> --previous`. The `--previous` flag is critical here because the current container is likely empty or restarting.
- Action: Fix the code bug or configuration syntax identified in the logs.



2.Exit Code 137 – OOMKilled

The "Why":

The container tried to utilize more memory than its limits.memory defined in the manifest. The Linux Out of Memory (OOM) Killer stepped in and sent a SIGKILL.

The Deep Dive:

Check the pod status using `kubectl describe pod`.
Look for Reason: OOMKilled.

The Fix:

- Vertical Scaling: Increase the memory limit in your YAML.
- Optimization: Investigate memory leaks or reduce the heap size (e.g., `-Xmx` for Java) to stay within the allocated "cgroup" boundaries.



3.ConfigMap/Secret Missing

The "Why":

The Pod is configured to load environment variables or mount volumes from a ConfigMap or Secret that does not exist or is misspelled.

The Deep Dive:

The Kubelet cannot even start the container. The Pod will stay in CreateContainerConfigError or jump straight to CrashLoopBackOff.

The Fix:

- Run `kubectl get configmap` and verify the name matches exactly (case-sensitive).
- Action: Create the missing resource or update the `envFrom / volumeMounts` section in your deployment.



4. Permission Denied (FSGroup)

The "Why":

The application process inside the container is trying to write to a mounted volume or a local directory (like /var/log) but doesn't have the Linux permissions.

The Deep Dive:

Common when running "Rootless" containers (SecurityContext runAsNonRoot: true).

The Fix:

- Use a securityContext in your YAML to set the fsGroup or runAsUser to a UID that owns that directory.
- Action: Adjust the Dockerfile CHOWN commands or Kubernetes securityContext.



5.Liveness Probe Failure

The "Why":

The application is actually running fine, but it's not responding to the "health check" (HTTP/TCP/Exec) defined in the manifest. Kubernetes thinks it's dead and kills it.

The Deep Dive:

Check `kubectl describe pod` for "Unhealthy" events. If the `initialDelaySeconds` is too short, K8s kills the app before it finishes booting.

The Fix:

- Increase `initialDelaySeconds` or `failureThreshold`.
- Action: Ensure the health check endpoint (e.g., `/healthz`) is actually implemented in your code.



6.Exit Code 127 – Command Not Found

The "Why":

The command or args specified in the Kubernetes manifest refers to a binary or script that doesn't exist inside the container image.

The Deep Dive:

This often happens when switching base images (e.g., moving from Ubuntu to Alpine, where bash is replaced by sh).

The Fix:

- Verify the ENTRYPOINT in your Dockerfile.
- Action: Use `docker run -it <image> /bin/sh` locally to verify the path of your executable.



7. Network/Dependency Timeout

The "Why":

The app starts, tries to connect to a Database or an external API, fails to connect within a timeout period, and performs a "Hard Exit."

The Deep Dive:

Classic "Microservice cascade." If the DB is down, the app crashes.

The Fix:

- Resilience: Implement a retry logic in your code so it doesn't crash on the first failed connection.
- Network: Check NetworkPolicies or Service names (DNS resolution). Use nslookup inside the cluster to verify connectivity.



8. Wrong Architecture (ARM vs AMD64)

The "Why":

You built the image on an Apple M-series chip (ARM64) and pushed it to a cloud cluster running Intel nodes (AMD64).

The Deep Dive:

The error in the logs will usually be:
`standard_init_linux.go:211: exec user process caused "exec format error".`

The Fix:

- Multi-arch builds: Use docker buildx to build for the correct target platform.
- Action: Re-push the image using the `--platform linux/amd64` flag.



9. Persistent Volume Mount Failure

The "Why":

The Pod is stuck because the PersistentVolumeClaim (PVC) it needs is already bound to a node in a different availability zone or is locked by another Pod (ReadWriteOnce).

The Deep Dive:

The container cannot start because its "disk" isn't available.

The Fix:

- Check `kubectl get pvc`. Ensure it is in Bound status.
- Action: Ensure your Nodes and Volumes are in the same AWS/GCP Zone. Use ReadWriteMany if multiple pods need the same disk.



10.Invalid Port Mapping

The "Why":

The application is listening on port 8080, but the Liveness Probe or Service is looking at port 80. The probe fails, causing a restart loop.

The Deep Dive:

This is a silent killer. Everything looks "green" in the logs, but the Pod keeps restarting every 30 seconds.

The Fix:

- Cross-reference the containerPort in the YAML with the actual port defined in your application code.
- Action: Align the ports in the livenessProbe and readinessProbe.



Bavithran
@bavicnative



Found this useful?



Follow



Let's connect



Found it useful? Drop your thoughts below and share it with your fellow DevOps engineers!

Follow me for more DevOps, Kubernetes, and cloud-native insights.



Repost